

# Advanced Architectures of Local Search: From Classical Operators to Neural-Symbolic Hybridization in Combinatorial Optimization

## 1. Introduction: The Landscape of Local Search in Computational Intractability

The resolution of NP-hard combinatorial optimization problems, particularly routing and scheduling variants such as the Traveling Salesman Problem (TSP) and the Vehicle Routing Problem (VRP), constitutes a cornerstone of modern operations research. Exact methods, while mathematically precise, suffer from exponential time complexity, rendering them computationally prohibitive for large-scale industrial instances involving thousands of nodes and complex constraints. Consequently, the burden of practical solvability falls upon heuristic and metaheuristic frameworks. At the heart of these frameworks lies the concept of **Local Search**—an iterative process of navigating the solution space by moving from one candidate solution to a neighboring solution through defined modifications known as **operators**.

The efficacy of any metaheuristic—be it Simulated Annealing, Tabu Search, or Genetic Algorithms—is fundamentally bounded by the quality and diversity of its underlying local search operators. These operators define the "neighborhood" topology; they determine which solutions are reachable from a given state and how easily the algorithm can escape local optima. Over the past six decades, the field has evolved from simple edge-exchange mechanisms designed for manual calculation to complex, variable-depth search procedures and, most recently, to heuristics generated automatically by deep neural networks and Large Language Models (LLMs).

This report provides an exhaustive analysis of this evolutionary trajectory. It dissects the mechanics, mathematical formulations, and computational implications of standard intra-route and inter-route operators. It explores the transition to Large Neighborhood Search (LNS) paradigms, details the mechanics of modern state-of-the-art (SOTA) approaches like Slack Induction by String Removal (SISR), and culminates in a critical review of the emerging field of Neural Combinatorial Optimization (NCO), where operators are learned rather than handcrafted.

## 2. Foundations of Intra-Route Optimization

Intra-route operators are the atomic units of local search. Originally conceived for the TSP,

their primary objective is to optimize the sequencing of nodes within a single tour to minimize travel cost or distance. These operators form the intensification phase of almost all complex VRP heuristics, refining the routes constructed by global algorithms.

## 2.1 The 2-opt Operator: Geometric Uncrossing and Path Inversion

The 2-opt operator, introduced by Croes in 1958 <sup>1</sup>, remains the most ubiquitous mechanism in routing optimization due to its elegance and profound effectiveness in Euclidean space.

### 2.1.1 The Mechanics of Inversion

Fundamentally, 2-opt is an edge-exchange operator. It operates by removing two non-adjacent edges from a tour and reconnecting the resulting segments to form a new valid tour. Consider a tour represented as a sequence of vertices  $T = \{v_1, v_2, \dots, v_n, v_1\}$ . If the operator selects edges  $(v_i, v_{i+1})$  and  $(v_j, v_{j+1})$  for removal (where  $i < j$ ), the only way to reconnect the path into a single Hamiltonian cycle—rather than two disjoint subtours—is to connect  $v_i$  to  $v_j$  and  $v_{i+1}$  to  $v_{j+1}$ .

This reconnection forces a topological inversion. The segment of the tour between  $v_{i+1}$  and  $v_j$  must be traversed in reverse order in the new solution. Thus, the new sequence becomes:

$$T' = \{v_1, \dots, v_i, v_j, v_{j-1}, \dots, v_{i+1}, v_{j+1}, \dots, v_n, v_1\}$$

This inversion property is critical. It implies that 2-opt is not merely swapping connections but physically untangling the route. In Euclidean instances, optimal tours generally do not contain crossing edges (due to the triangle inequality). The 2-opt operator is the primary mechanism for eliminating such crossings, as replacing crossing edges with non-crossing ones almost always reduces the tour length.<sup>1</sup>

### 2.1.2 Algorithmic Implementation and Complexity

A naive implementation of 2-opt might re-calculate the total length of the tour for every possible pair of edges, leading to an  $O(n)$  cost per move evaluation and an  $O(n^3)$  complexity for a full neighborhood scan. However, efficient implementations utilize differential evaluation (or delta evaluation).

The cost change  $\Delta$  for a 2-opt move involving edges  $(v_i, v_{i+1})$  and  $(v_j, v_{j+1})$  is calculated as:

$$\Delta = d(v_i, v_j) + d(v_{i+1}, v_{j+1}) - d(v_i, v_{i+1}) - d(v_j, v_j)$$

where  $d(u, v)$  denotes the distance between vertices  $u$  and  $v$ .

Because the internal edges of the reversed segment  $(v_{i+1}, \dots, v_j)$  remain identical (only their direction changes), their contribution to the total cost remains constant in symmetric TSP (where  $d(a, b) = d(b, a)$ ). This reduces the evaluation of a move to an  $O(1)$  operation involving only four distance lookups. Consequently, a full scan of the 2-opt neighborhood requires  $O(n^2)$  time. In asymmetric TSP, where  $d(a, b) \neq d(b, a)$ , the cost of the reversed segment changes, necessitating either an  $O(n)$  evaluation or complex auxiliary data structures to maintain segment costs.<sup>1</sup>

## 2.2 The 3-opt Operator and Permutation Complexity

While 2-opt is effective, it can get trapped in local optima where no 2-edge exchange yields an improvement, yet the solution is not globally optimal. The 3-opt operator generalizes the concept by removing three edges, creating three path segments.

### 2.2.1 Reconnection Topologies

Removing three edges allows for a richer set of reconnections. There are theoretically  $(2^3 \times 3!)/(2 \times 3) = 8$  ways to reconnect three segments to form a tour, one of which is the original tour. The remaining seven moves constitute the 3-opt neighborhood. These can be categorized into:

1. **Pure 3-opt moves:** These involve reconnections that cannot be achieved by sequential 2-opt moves. They allow for the relocation of a segment without reversing it, or reversing two separate segments.
2. **Degenerate moves:** Some 3-opt configurations are equivalent to performing two or three 2-opt moves in sequence.

The complexity of a full 3-opt neighborhood scan is  $O(n^3)$ , which is often prohibitively expensive for large instances ( $n > 1000$ ). Therefore, 3-opt is frequently implemented using "candidate lists" (checking only the  $k$ -nearest neighbors) or as a "3-opt perturbation" in Iterated Local Search rather than an exhaustive search operator.<sup>4</sup>

### 2.2.2 The Case of the 3-Permute Operator

Recent investigations into specific variants, such as the **3-permute** operator, highlight the nuances of operator design. In a study comparing operators for Open-Loop TSP (where the tour is a path, not a cycle), the 3-permute operator was analyzed for its ability to reorder segments. However, experimental results revealed that 3-permute failed to produce unique results distinct from combinations of Relocate and 2-opt. In open-loop scenarios, the structural advantage of 3-permute vanished, and it was dominated by operators that specifically manipulated the endpoints of the path, such as the **Link Swap** operator.<sup>4</sup>

## 2.3 Node-Centric Operators: Relocate and Swap

While 2-opt and 3-opt are edge-centric (focusing on connections), Relocate and Swap are node-centric (focusing on the position of customers).

### 2.3.1 Relocate (Or-opt)

The Relocate operator (also known as Shift or Insert) removes a vertex  $v_i$  from its current position and reinserts it at a new position  $j$ .

- **Mechanism:** Delete  $(v_{i-1}, v_i)$  and  $(v_i, v_{i+1})$ . Add  $(v_{i-1}, v_{i+1})$ . Insert  $v_i$  between  $v_j$  and  $v_{j+1}$  by adding  $(v_j, v_i)$  and  $(v_i, v_{j+1})$  and removing  $(v_j, v_{j+1})$ .
- **Or-opt:** A generalization proposed by Or (1976) involves relocating a *chain* of consecutive vertices (typically of length 1, 2, or 3). This is particularly effective for VRPs where a sequence of customers might be geographically clustered and should be moved together.
- **Utility:** Relocate is essential for "fine-tuning." While 2-opt fixes large structural crossings, Relocate handles local sequencing errors where a customer is visited slightly too early or too late in the path.<sup>6</sup>

### 2.3.2 Exchange (Swap)

The Swap operator selects two vertices  $v_i$  and  $v_j$  and exchanges their positions.

- **Mechanism:** Unlike 2-opt, Swap does *not* invert the path segment between the two nodes. It preserves the orientation of the intermediate chain.
- **Use Case:** This operator is vital when the general direction of flow is correct (so inversion would be detrimental) but specific assignments are suboptimal. For example, if a vehicle passes two clusters but picks the wrong representative from each, Swap can correct this without disrupting the intra-cluster paths.<sup>6</sup>

## 3. Inter-Route Operators: The Mechanics of VRP Optimization

In Vehicle Routing Problems, the optimization challenge expands from sequencing (TSP) to clustering and assigning customers to different vehicles. Intra-route operators cannot move a customer from one vehicle to another; thus, they cannot balance loads or reduce the fleet size. Inter-route operators are designed to facilitate these exchanges between routes.

### 3.1 Cross-Exchange and the $\lambda$ -Interchange Framework

The Cross-Exchange operator is a powerful mechanism that generalizes simply moving a node. It involves swapping two sequences of vertices between two different routes, Route A and Route B.

#### 3.1.1 The Mechanics of Segment Exchange

Let Route A contain a segment of customers  $S_A$  and Route B contain a segment  $S_B$ . The Cross-Exchange operator swaps  $S_A$  with  $S_B$ .

- **Symmetry:** Unlike Relocate (which is unidirectional), Cross-Exchange is bidirectional.
- **Preservation of Orientation:** Crucially, the internal order of customers within  $S_A$  and  $S_B$  is preserved. This is highly advantageous for VRP with Time Windows (VRPTW), as a valid sequence of customers is likely to remain valid (locally) when moved to a new route, provided the arrival time at the insertion point is feasible.

#### 3.1.2 $\lambda$ -Interchange Generalization

Osman (1993) formalized this behavior in the  $\lambda$ -Interchange framework.<sup>8</sup> The parameter  $\lambda$  denotes the maximum size of the segment exchanged.

- If  $|S_A| = 1$  and  $|S_B| = 0$ , the move is a **Relocate (Shift)**  $(1, 0)$ .
- If  $|S_A| = 1$  and  $|S_B| = 1$ , the move is a **Swap**  $(1, 1)$ .
- If  $|S_A| = 2$  and  $|S_B| = 2$ , the move is a **Swap(2,2)**.

The computational complexity of a full  $\lambda$ -Interchange neighborhood is  $O(n^2 \lambda^2)$ . By limiting  $\lambda$  to small values (e.g., 2 or 3), the operator remains efficient while providing the ability to move clusters of nodes, which is essential for escaping local optima where single-node moves are blocked by capacity constraints.<sup>8</sup>

#### 3.1.3 Delta Evaluation for Cross-Exchange

The delta evaluation for Cross-Exchange is  $O(1)$  assuming fixed  $\lambda$ . The cost change is

derived solely from the breaking of links at the start and end of the segments and the formation of new links. No internal edges of  $S_A$  or  $S_B$  are modified.

$$\Delta = d(u_{pre}, v_{start}) + d(u_{end}, v_{post}) + d(v_{pre}, u_{start}) +$$

where  $u$  nodes belong to the segment from Route A and  $v$  nodes to Route B. The complexity arises not from the cost calculation, but from the feasibility checks (capacity and time windows) for the altered routes.<sup>1</sup>

### 3.2 2-opt\* (2-opt Star): The Tail Swapper

Standard 2-opt, if applied to two different routes, would effectively merge them into one giant route (by deleting one edge from each and cross-connecting). In VRP, this often violates vehicle capacity. The 2-opt\* operator, proposed by Potvin and Rousseau (1995), is a modification specifically for inter-route optimization.

#### 3.2.1 Mechanism

2-opt\* selects two routes and an edge from each:  $(u_1, u_2)$  from Route A and  $(v_1, v_2)$  from Route B. It deletes these edges and creates new connections  $(u_1, v_2)$  and  $(v_1, u_2)$ .

- **Effect:** This effectively swaps the "tails" of the two routes. All customers after  $u_1$  are moved to Route B, and all customers after  $v_1$  are moved to Route A.
- **Feasibility:** Unlike standard 2-opt, 2-opt\* preserves the orientation of the segments. This makes it particularly robust for Time Window constraints, as the relative order of customers in the tails is not reversed.<sup>6</sup>

### 3.3 Swap\* (Swap Star): Optimization of Insertion

The standard Swap operator exchanges two nodes  $u$  and  $v$  in their respective positions.

However, the best position for  $u$  in Route B might not be exactly where  $v$  was sitting. The Swap\* operator, introduced by Vidal et al., addresses this.

- **Mechanism:** Swap\* removes  $u$  from Route A and  $v$  from Route B. It then calculates the *optimal insertion cost* for  $u$  into any position in Route B and  $v$  into any position in Route A.
- **Complexity:** A naive approach would be  $O(n^2)$  for the positions  $\times O(n^2)$  for the

pairs, leading to  $O(n^4)$ . However, efficient implementations use pre-calculated insertion costs and pruning to achieve sub-quadratic performance.

- **Impact:** Swap\* is significantly more powerful than simple Swap for fleet size reduction and load balancing, as it allows nodes to migrate to their "natural" positions in the target routes rather than being forced into arbitrary slots.<sup>8</sup>

## 4. Advanced Neighborhoods and Variable Depth Search

When simple 1-move or 2-move neighborhoods are exhausted (local optima), advanced heuristics employ variable depth searches or complex compound moves to find improvements.

### 4.1 Ejection Chains: The Domino Effect

Ejection Chains, pioneered by Glover (1996) and Rego, operate on the principle that a move which is locally infeasible (e.g., inserting a node into a full vehicle) can be made feasible by "ejecting" another element, which in turn displaces another, creating a chain reaction.

#### 4.1.1 Mechanics and Reference Structures

An ejection chain is defined by a sequence of moves where the feasibility of step  $k$  depends on the execution of step  $k + 1$ . The chain continues until a "closing" condition is met, such as inserting the final ejected element into an empty slot or the slot originally vacated by the first element.

- **Reference Structures:** To manage the complexity of possible chains, algorithms use "reference structures."
  - **Stem-and-Cycle:** A common structure where a path (stem) leads to a cycle of ejections.
  - **Flower Structure:** A more complex topology used in VRP to explore multi-route exchanges simultaneously.<sup>11</sup>
- **Application:** Ejection chains are the primary mechanism for **Fleet Size Minimization**. To reduce the number of vehicles, the algorithm attempts to empty a route by ejecting its customers one by one into other routes. If a target route is full, an ejection chain is triggered to make space.<sup>12</sup>

### 4.2 The Link Swap Operator: Dominance in Open-Loop TSP

In 2019, research specifically targeting Open-Loop TSP (Path TSP) revealed the limitations of standard 2-opt. In a closed loop, every vertex has degree 2. In an open loop, the start and end vertices have degree 1, and standard 2-opt (which assumes a cycle) cannot easily manipulate

these endpoints without creating invalid topologies or requiring virtual "dummy" edges.

#### 4.2.1 The Link Swap Alternative

The Link Swap operator was developed to address this. For a selected link (edge), it generates three neighbors:

1. **Alternative 1:** Equivalent to a standard 2-opt.
2. **Alternative 2:** Equivalent to a standard 2-opt.
3. **Alternative 3 (The Endpoint Switch):** This unique move connects the two previous *terminal* points of the path.

In closed-loop TSP, this third alternative is redundant or invalid. However, in open-loop scenarios, it allows the path to "flip" inside out, making the old endpoints internal and promoting two internal nodes to endpoints. Experimental analysis showed that while 2-opt and Relocate share improvements equally in closed loops, **Link Swap accounts for 50% of all improvements** in open-loop instances, effectively rendering it the SOTA operator for that specific variant.<sup>4</sup>

### 4.3 Lin-Kernighan-Helsgaun (LKH)

No discussion of local search is complete without mentioning the Lin-Kernighan (LK) heuristic.

LK is effectively a dynamic  $k$ -opt algorithm. It decides the value of  $k$  adaptively during the search. It builds a sequence of 2-opt moves, locking edges that have been added to prevent cycling, and backtracks if the chain does not lead to an improvement. The Helsgaun implementation (LKH) includes highly sophisticated candidate set management (using Minimum Spanning Trees and Alpha-closeness) and remains the undefeated champion for solving large-scale symmetric TSPs.<sup>14</sup>

## 5. Large Neighborhood Search (LNS) and Adaptive Frameworks

As problem complexity increases (e.g., Rich VRP with multiple constraints), local search neighborhoods become riddled with infeasible valleys. Small moves like 2-opt struggle to traverse these valleys. **Large Neighborhood Search (LNS)**, and its adaptive variant **ALNS** (Ropke & Pisinger, 2006), shift the paradigm from "moving" elements to "destroying and repairing" the solution.

### 5.1 The Ruin (Destroy) Phase

The core philosophy of LNS is to remove a significant portion of the solution (e.g., 10-40% of customers) to create a partial solution. The method of removal dictates the trajectory of the search.

### 5.1.1 Random Removal

Simply removes  $q$  customers chosen uniformly at random. This serves as a diversification mechanism (a "kick" to the system) but rarely leads to intensification or finding high-quality local optima on its own.<sup>16</sup>

### 5.1.2 Worst Removal

This operator introduces a greedy bias. It removes customers that are "expensive" in the current configuration. The cost of a customer  $c$  is typically defined as  $Cost(S) - Cost(S \setminus \{c\})$ .

- **Stochasticity:** To avoid deterministic looping, a randomization factor is added. The customers are sorted by their removal "savings," and the algorithm selects the  $i$ -th worst customer where  $i$  is determined by a randomized power law ( $y^p$ ).<sup>16</sup>

### 5.1.3 Cluster Removal

This operator targets geographic decomposition.

- **Mechanism:** It selects a seed customer and removes a cluster of neighbors around it. This is often implemented using a modified **Kruskal's algorithm** (Minimum Spanning Tree) or simply distance-based radial selection.
- **Rationale:** Removing a geographic cluster creates a large "hole" in the map. This allows the repair operator to fundamentally restructure the service logic for that region—perhaps servicing it with a different vehicle or entering/exiting the cluster from different angles—something 2-opt cannot easily achieve.<sup>19</sup>

### 5.1.4 Shaw Removal (Related Removal)

Proposed by Shaw (1998), this operator is based on the idea that it is easier to swap customers that are "related" or similar.

- **Relatedness Metric:** A metric  $R(i, j)$  computes the similarity between customer  $i$  and  $j$  based on:
  1. **Distance:**  $d(i, j)$
  2. **Time:**  $|T_i - T_j|$  (Difference in service windows)
  3. **Demand:**  $|q_i - q_j|$
  4. **Vehicle compatibility.**
- **Execution:** The operator selects a seed, then iteratively removes customers that are most "related" to the ones already removed. This effectively removes a set of customers

that could potentially be served by the same vehicle or swapped with each other, maximizing the rearrangement potential during repair.<sup>18</sup>

## 5.2 The Repair (Recreate) Phase

Once the solution is destroyed, it must be rebuilt. The quality of the LNS step depends heavily on the intelligence of the insertion.

### 5.2.1 Regret Insertion Heuristics

A simple Greedy Insertion places the customer with the lowest insertion cost first. However, this is myopic; placing customer A cheaply now might force customer B to be inserted later at a massive cost. Regret heuristics mitigate this look-ahead problem.

- **Regret-2:** For each unassigned customer  $u$ , calculate:
  - $C_1(u)$ : Cost of inserting  $u$  in its best route.
  - $C_2(u)$ : Cost of inserting  $u$  in its second best route.
  - **Regret Value:**  $\Delta(u) = C_2(u) - C_1(u)$ .
  - **Selection:** Insert the customer with the **maximum** regret value.
  - **Logic:** "We must insert this customer now, because if we wait and miss its best route, the alternative is much more expensive."
- **Regret-k:** Generalizes this to the difference between the best and the  $k$ -th best route. Regret-3 and Regret-4 are common variants that look deeper into the route options. Empirical studies show that Regret-2 and Regret-3 generally outperform basic greedy insertion in almost all VRP variants.<sup>22</sup>

## 6. State-of-the-Art (2020-2024): SISR and Slack Induction

In the domain of handcrafted heuristics, the **Slack Induction by String Removal (SISR)** algorithm, introduced by Christiaens and Vanden Berghe (2020), represents the current state-of-the-art for a wide range of VRP benchmarks. It is notable for outperforming complex ALNS implementations while being conceptually simpler.

### 6.1 The Concept of Spatial Slack

Traditional VRP heuristics focus on "Capacity Slack" (how much room is left in the truck). SISR introduces the concept of **Spatial Slack**.

- **Definition:** Spatial slack refers to the detour flexibility of a route. If a route is extremely tight (a straight line between customers), any insertion incurs high cost.
- **Visualizing Slack:** Imagine an ellipse around each edge of the route, with the endpoints

as foci. The size of the ellipse represents the area reachable within a certain cost bound.

- **Induction:** When a heuristic removes nodes, it creates gaps. SISR specifically aims to induce *spatial* slack by removing **adjacent strings** of customers. Removing a sequence of nodes (rather than scattered random nodes) creates a large, contiguous amount of maneuverability (slack) in the route, allowing it to "bend" significantly to accept new, distant insertions.<sup>26</sup>

## 6.2 Ruin Strategy: Adjacent String Removal

SISR deviates from Shaw or Worst removal.

- **String-Based:** It selects a random route and a random seed customer, then removes a contiguous string of length  $L$ .
- **Neighborhood Preservation:** It then looks at the *neighbors* of the removed string (using a precomputed distance matrix) and removes strings from *their* routes as well.
- **Result:** This creates a geographically concentrated "disaster zone" across multiple routes, dissolving the local structure entirely in that region. This prevents the repair operator from simply falling back into the previous local optimum.<sup>26</sup>

## 6.3 Recreate Strategy: Greedy Insertion with Blinks

For the repair phase, SISR uses a standard Greedy Insertion but with a crucial modification called **Blinks**.

- **The Cost of Repair:** Standard insertion checks every customer against every position, which is  $O(n^2)$ .
- **Blinking:** To allow for massive iteration counts (millions of moves), SISR's insertion operator "blinks" (skips) a check of valid insertion positions with a probability  $p$ .
- **Effect:** This acts as a randomized approximation. It speeds up the repair drastically, allowing the algorithm to explore the search space much faster. The combination of **Spatial Slack Induction** (via string removal) and **High-Frequency Randomized Repair** (via blinks) allows SISR to converge to solutions superior to those found by complex ALNS frameworks.<sup>27</sup>

# 7. State-of-the-Art (2024-2025): Neural Combinatorial Optimization (NCO)

The frontier of local search is shifting from "designing operators" to "learning operators" using Deep Reinforcement Learning (DRL) and Graph Neural Networks (GNNs). This field, known as Neural Combinatorial Optimization, aims to replace hand-coded heuristics with learned policies.

## 7.1 Learning to Delegate (Sub-problem Selection)

One major bottleneck in LNS is deciding *which* part of the solution to destroy.

- **GNN-Based Heatmaps:** Recent approaches (e.g., "Learning to Delegate" by Li et al.) use GNNs to predict a "heatmap" of the solution edges. The network assigns probabilities to edges indicating their likelihood of being part of the optimal solution.
- **Neural Destroy:** The LNS destroy operator uses this heatmap to remove edges with low probability, effectively learning a "Worst Removal" operator that is topologically aware rather than just cost-aware.<sup>30</sup>

## 7.2 Learning Operator Selection

Instead of using a roulette wheel (as in ALNS) to select between 2-opt, Relocate, and Swap, RL agents are trained to select the best operator for the current state.

- **State Representation:** The state is encoded as a graph (node coordinates, demands, current routes).
- **Policy Network:** An RL agent (often PPO or DQN) outputs a discrete action: "Apply 2-opt on Route 5".
- **Reward:** The reward is the improvement in solution cost.
- **Advantage:** The neural network learns context. It might learn that "when routes are crossing near the depot, 2-opt\* is best," but "when a route is zig-zagging far away, Relocate is best"—rules that are difficult to hard-code manually.<sup>32</sup>

## 7.3 Neural Repair Operators

Deep Learning models are also being used to perform the repair step.

- **Generative Reconstruction:** Sequence-to-sequence models (like Transformers) can take the partial solution (after destruction) and autoregressively generate the missing sequence of customers, effectively learning a complex insertion heuristic that considers global context rather than just local regret.<sup>30</sup>

# 8. State-of-the-Art (2025): LLM-Driven Heuristic Discovery

The most radical innovation in 2025 is the use of Large Language Models (LLMs) not just to *run* heuristics, but to *write* them.

## 8.1 VRPAGent and Evolutionary Code Generation

**VRPAGent** (2025) represents a paradigm shift. It integrates an LLM into an evolutionary search framework.<sup>34</sup>

- **LLM as Mutation Operator:** Instead of mutating numerical parameters, the algorithm

mutates Python code. The LLM is prompted with the code of a current heuristic (e.g., a simple greedy removal) and asked to "improve this heuristic to prioritize customers with tight time windows."

- **Genetic Cycle:**
  1. **Generate:** The LLM writes 10 variations of Python code for a destroy operator.
  2. **Evaluate:** These scripts are executed on a set of VRP instances.
  3. **Select:** The best-performing code snippets are kept.
  4. **Refine:** The best code is fed back to the LLM with performance stats and a prompt to "optimize further."
- **Outcome:** This system has discovered novel operators that human experts had not formalized, such as hybrid removal operators that combine spatial density with time-window urgency in non-linear formulas. It automates the role of the algorithm designer.<sup>36</sup>

## 9. Implementation Mechanics and Computational Complexity

Theoretical elegance is useless without computational efficiency. The viability of any local search operator hinges on its ability to be evaluated in  $O(1)$  or  $O(n)$  time.

### 9.1 Delta Evaluation Formulas

The key to efficiency is  $\Delta$ -evaluation: calculating the cost difference without traversing the whole route.

Operator	Complexity	Delta Formula Concept
2-opt	$O(1)$	$d(u, v') + d(v, u') - d(u, u')$
Relocate	$O(1)$	Remove 2 edges, Add 3 edges (triangular reconnection).
Cross-Exchange	$O(1)$	Depends only on the 4 cut-points (start/end of both segments). Internal edges cancel out.
2-opt*	$O(1)$	Connect tail of A to head of

		B and vice versa.
3-opt	$O(1)$	3 removed edges, 3 added. 7 cases of reconnection.

**Note on Asymmetry:** In Asymmetric TSP (ATSP) or VRP with time-dependent travel times, path reversal (in 2-opt) changes internal edge costs. In these cases, 2-opt becomes  $O(n)$  unless complex segment-tree data structures are used.<sup>1</sup>

## 9.2 Data Structures for Speed

- **Doubly Linked Lists:** Standard arrays are inefficient for insertion/deletion ( $O(n)$  shift). Linked lists allow  $O(1)$  structural changes, but make delta-evaluation harder (no random access).
- **Two-Level Tree / Segment Tree:** A hybrid structure where the route is broken into segments. This allows  $O(\log n)$  operations for both testing and execution, supporting operations like "Flip Segment" efficiently even in asymmetric contexts.
- **Candidate Lists:** For large graphs, checking every pair of edges for 2-opt ( $O(n^2)$ ) is too slow. A **Candidate List** stores the  $k$  nearest neighbors for each node. The search only considers moves that introduce an edge to a neighbor in the list, reducing complexity to  $O(kn)$ .<sup>14</sup>
- **Don't Look Bits (DLB):** A simple but vital optimization. A bit is assigned to each node. If a search centered on Node A fails to find an improvement, the bit is set to 1. In the next iteration, Node A is skipped. The bit is reset to 0 only if one of Node A's neighbors is modified. This drastically reduces wasted CPU cycles.<sup>15</sup>

## 9.3 GPU Parallelization

With the rise of CUDA-accelerated solvers (e.g., NVIDIA cuOpt), local search is being parallelized.

- **Batch Move Evaluation:** Instead of evaluating one move at a time, thousands of potential 2-opt moves are evaluated in parallel threads.
- **Independent Route Processing:** In LNS, the repair of Route A and Route B can happen simultaneously on different GPU blocks if they are spatially disjoint.<sup>37</sup>

# 10. Conclusion

The domain of local search operators has undergone a profound transformation. While the

geometric foundations laid by **2-opt** (1958) and **3-opt** (1965) remain the bedrock of intra-route optimization, the field has ascended to higher levels of abstraction to tackle the complexity of modern logistics.

The transition from simple moves to **Large Neighborhood Search** marked the first major leap, acknowledging that "destruction" is as important as "construction." The development of **SISR (2020)** refined this further by introducing **Spatial Slack**, demonstrating that topological flexibility is the key to escaping local optima in tightly constrained problems.

However, the current era (2024-2025) is defined by the **hybridization of symbolic and neural methods**. The emergence of **Neural Combinatorial Optimization** and **LLM-driven VRP Agent** indicates a future where operators are no longer static tools selected from a library, but dynamic, learned policies generated on the fly. These systems do not just solve the routing problem; they solve the "algorithm design problem," tailoring the local search operators themselves to the specific distribution of the data at hand.

**Table 1: Summary of Operator Classifications and Characteristics**

Operator Class	Operator Name	Key Mechanism	Best For	Complexity (Scan)
Intra-Route	2-opt	Edge crossing elimination	Euclidean TSP/VRP	$O(n^2)$ / $O(kn)$
Intra-Route	Or-opt (Relocate)	Moving chains of nodes	Fine-tuning sequences	$O(n^2)$
Inter-Route	Cross-Exchange	Swapping segments	VRP with Time Windows	$O(n^2 \lambda^2)$
Inter-Route	2-opt*	Swapping tails	VRP Fleet reduction	$O(n^2)$
Inter-Route	Swap*	Optimal re-insertion	VRP Load Balancing	Sub- $O(n^2)$
Meta-Operator	Ejection Chain	Compound displacement	Constrained VRP	Variable
LNS Destroy	Cluster	Geographic	Escaping local	$O(n \log n)$

	Removal	hole creation	optima	
<b>LNS Repair</b>	Regret-k	Look-ahead insertion	High-quality reconstruction	$O(n^2)$
<b>SOTA Heuristic</b>	SISR	Spatial Slack + Blinks	Rich VRP / Benchmarks	Linear-time phases
<b>SOTA Neural</b>	Neural Destroy	Heatmap-based removal	Learned distributions	Inference time
<b>Specialized</b>	Link Swap	Endpoint switching	Open-Loop TSP	$O(1)$ per link

## Works cited

1. 2-opt - Wikipedia, accessed February 4, 2026, <https://en.wikipedia.org/wiki/2-opt>
2. A Strategy for Reducing the Computational Complexity of Local Search-Based Methods, and its Application to the Vehicle Routing Problem, accessed February 4, 2026, [http://users.ntua.gr/ezach/A\\_Strategy\\_for\\_Reducing\\_the\\_Computational\\_Complexity\\_of\\_Local\\_Search-Based\\_Methods\\_and\\_its\\_Application\\_to\\_the\\_VRP.pdf](http://users.ntua.gr/ezach/A_Strategy_for_Reducing_the_Computational_Complexity_of_Local_Search-Based_Methods_and_its_Application_to_the_VRP.pdf)
3. How to Improve TSP-Tours Applying the 2-opt Neighborhood - phabe.ch, accessed February 4, 2026, <https://phabe.ch/2024/08/27/how-to-improve-tsp-tours-applying-the-2-opt-neighborhood/>
4. Which Local Search Operator Works Best for the Open-Loop TSP? - MDPI, accessed February 4, 2026, <https://www.mdpi.com/2076-3417/9/19/3985>
5. How does the 3-opt algorithm for TSP work? - Computer Science Stack Exchange, accessed February 4, 2026, <https://cs.stackexchange.com/questions/19808/how-does-the-3-opt-algorithm-for-tsp-work>
6. 2-opt – Knowledge and References - Taylor & Francis, accessed February 4, 2026, [https://taylorandfrancis.com/knowledge/Engineering\\_and\\_technology/Computer\\_science/2-opt/](https://taylorandfrancis.com/knowledge/Engineering_and_technology/Computer_science/2-opt/)
7. Generalized vehicle routing problem: Contemporary trends and research directions - PMC, accessed February 4, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC10731084/>
8. Heuristics for Vehicle Routing Problem: A Survey and Recent Advances - arXiv, accessed February 4, 2026, <https://arxiv.org/pdf/2303.04147.pdf>
9. Andelmin, Juho; Bartolini, Enrico A multi-start local search heuristic for the green vehicle routing problem based on a multigra, accessed February 4, 2026,

[https://research.aalto.fi/files/33842218/SCI\\_Andeelman\\_Bartolini\\_A\\_Multi\\_Start\\_Local\\_Search.pdf](https://research.aalto.fi/files/33842218/SCI_Andeelman_Bartolini_A_Multi_Start_Local_Search.pdf)

10. O(1) Delta Component Computation Technique for the Quadratic Assignment Problem - ResearchGate, accessed February 4, 2026,  
[https://www.researchgate.net/publication/225098587\\_O1\\_Delta\\_Component\\_Computation\\_Technique\\_for\\_the\\_Quadratic\\_Assignment\\_Problem/fulltext/030478f60cf262df8c19b5aa/O1-Delta-Component-Computation-Technique-for-the-Quadratic-Assignment-Problem.pdf](https://www.researchgate.net/publication/225098587_O1_Delta_Component_Computation_Technique_for_the_Quadratic_Assignment_Problem/fulltext/030478f60cf262df8c19b5aa/O1-Delta-Component-Computation-Technique-for-the-Quadratic-Assignment-Problem.pdf)
11. An ejection chain algorithm for the quadratic assignment problem - University of Colorado Boulder, accessed February 4, 2026,  
<https://leeds-faculty.colorado.edu/glover/fred%20pubs/401%20-%20Ejection%20Chain%20for%20QAP%20-%20w%20Cesar%20&%20Tabitha%20-%20Networks.pdf>
12. Fast Ejection Chain Algorithms for Vehicle Routing with Time Windows \* - Marc Uetz, accessed February 4, 2026,  
[https://marcuetz.personalweb.utwente.nl/Preprints/ejection\\_chain\\_paper.pdf](https://marcuetz.personalweb.utwente.nl/Preprints/ejection_chain_paper.pdf)
13. Fast Local Searches For The Vehicle Routing Problem With Time Windows - ResearchGate, accessed February 4, 2026,  
[https://www.researchgate.net/publication/221704664\\_Fast\\_Local\\_Searches\\_For\\_The\\_Vehicle\\_Routing\\_Problem\\_With\\_Time\\_Windows](https://www.researchgate.net/publication/221704664_Fast_Local_Searches_For_The_Vehicle_Routing_Problem_With_Time_Windows)
14. The Traveling Salesman Problem: State of the Art - UBC Computer Science, accessed February 4, 2026, <https://www.cs.ubc.ca/~hoos/SLS-Book/Slides/tsp.pdf>
15. Travelling Salesman Problems - UBC Computer Science, accessed February 4, 2026,  
<https://www.cs.ubc.ca/labs/algorithms/Courses/CPSC532D-05/Slides/tsp-camilo.pdf>
16. Large Multiple Neighborhood Search for the Clustered Vehicle-Routing Problem, accessed February 4, 2026,  
<https://logistik.bwl.uni-mainz.de/files/2018/12/LM-2017-01.pdf>
17. An Adaptive Large Neighborhood Search for the Larger-Scale Instances of Green Vehicle Routing Problem with Time Windows - ResearchGate, accessed February 4, 2026,  
[https://www.researchgate.net/publication/346576408\\_An\\_Adaptive\\_Large\\_Neighborhood\\_Search\\_for\\_the\\_Larger-Scale\\_Instances\\_of\\_Green\\_Vehicle\\_Routing\\_Problem\\_with\\_Time\\_Windows](https://www.researchgate.net/publication/346576408_An_Adaptive_Large_Neighborhood_Search_for_the_Larger-Scale_Instances_of_Green_Vehicle_Routing_Problem_with_Time_Windows)
18. Application of Adaptive Large Neighbourhood Search for a Rich and Real-World Vehicle Routing Problem - ResearchGate, accessed February 4, 2026,  
[https://www.researchgate.net/profile/Mohamed\\_Mourad\\_Lafifi/post/Is\\_SISR\\_S\\_Lack\\_Induction\\_by\\_String\\_Removals\\_for\\_Vehicle\\_Routing\\_Problems\\_successfully\\_implemented\\_anywhere/attachment/5e16c75ecfe4a777d4025aa6/AS%3A845318878932992%401578551134859/download/Application+of+Adaptive+Large+Neighborhood+Search+for+a+Rich+and+Real-World+Vehicle+Routing+Problem++Krop\\_MA\\_BMS.pdf](https://www.researchgate.net/profile/Mohamed_Mourad_Lafifi/post/Is_SISR_S_Lack_Induction_by_String_Removals_for_Vehicle_Routing_Problems_successfully_implemented_anywhere/attachment/5e16c75ecfe4a777d4025aa6/AS%3A845318878932992%401578551134859/download/Application+of+Adaptive+Large+Neighborhood+Search+for+a+Rich+and+Real-World+Vehicle+Routing+Problem++Krop_MA_BMS.pdf)
19. Adaptive Large Neighborhood Search, accessed February 4, 2026,  
<https://d-nb.info/1072464683/34>

20. A Hybrid Brain Storm Optimization Algorithm to Solve the Emergency Relief Routing Model, accessed February 4, 2026,  
<https://www.mdpi.com/2071-1050/15/10/8187>
21. Large neighborhood search for multi-trip vehicle routing - ORBi, accessed February 4, 2026,  
<https://orbi.uliege.be/bitstream/2268/187309/1/LNS%20for%20the%20MTVRP%20-%20WP%2020151130.pdf>
22. Greedy Heuristics with Regret, with Application to the Cheapest Insertion Algorithm for the TSP, accessed February 4, 2026,  
[https://www.tau.ac.il/~hassin/tsp\\_regret.pdf](https://www.tau.ac.il/~hassin/tsp_regret.pdf)
23. Heuristic Algorithms - Master's Degree in Computer Science/Mathematics, accessed February 4, 2026,  
<https://homes.di.unimi.it/cordone/courses/2024-ae/Lez09.pdf>
24. An Adaptive Large Neighborhood Search Heuristic for a Multi-Period Vehicle Routing Problem - Cirreli, accessed February 4, 2026,  
<https://www.cirreli.ca/documents/travail/cirreli-2013-67.pdf>
25. An adaptive large neighborhood search heuristic for Two-Echelon Vehicle Routing Problems arising in city logistics - PMC, accessed February 4, 2026,  
<https://pmc.ncbi.nlm.nih.gov/articles/PMC3587400/>
26. Slack Induction by String Removals for Vehicle Routing Problems - ResearchGate, accessed February 4, 2026,  
[https://www.researchgate.net/publication/338614875\\_Slack\\_Induction\\_by\\_String\\_Removals\\_for\\_Vehicle\\_Routing\\_Problems](https://www.researchgate.net/publication/338614875_Slack_Induction_by_String_Removals_for_Vehicle_Routing_Problems)
27. SISR Slack Induction by String Removals | PDF | Mathematical Optimization - Scribd, accessed February 4, 2026,  
<https://www.scribd.com/document/797815670/SISR-Slack-induction-by-string-removals>
28. Development and investigation of a sequence-based hyper-heuristic for vehicle routing problems - University of Exeter, accessed February 4, 2026,  
<https://ore.exeter.ac.uk/n downloader/files/58174189>
29. Slack Induction by String Removals for Vehicle Routing Problems - IDEAS/RePEc, accessed February 4, 2026,  
<https://ideas.repec.org/a/inm/ortrsc/v54y2020i2p417-433.html>
30. VRPAgent: LLM-Driven Discovery of Heuristic Operators for Vehicle Routing Problems, accessed February 4, 2026,  
<https://www.semanticscholar.org/paper/VRPAgent%3A-LLM-Driven-Discovery-of-Heuristic-for-Hotung-Berto/1d1a55403692e98cf936f16e957672f95d400011>
31. Neural Combinatorial Optimization Algorithms for Solving Vehicle Routing Problems: A Comprehensive Survey with Perspectives - arXiv, accessed February 4, 2026, <https://arxiv.org/html/2406.00415v1>
32. Learning local search operator selection - Operations Research Stack Exchange, accessed February 4, 2026,  
<https://or.stackexchange.com/questions/5119/learning-local-search-operator-selection>
33. Automated Design of Search Algorithms: Learning on Algorithmic Components -

- University of Nottingham, accessed February 4, 2026,  
<https://people.cs.nott.ac.uk/pszrq/files/ESWA21.pdf>
- 34. [2510.07073] VRPAGENT: LLM-Driven Discovery of Heuristic Operators for Vehicle Routing Problems - arXiv, accessed February 4, 2026,  
<https://arxiv.org/abs/2510.07073>
  - 35. VRPAGENT: LLM-DRIVEN DISCOVERY OF HEURISTIC OPERATORS FOR VEHICLE ROUTING PROBLEMS - OpenReview, accessed February 4, 2026,  
<https://openreview.net/pdf/bde8d2c4bfb0cd193ab13b7ab73b99416d623e50.pdf>
  - 36. VRPAGENT: LLM-Driven Discovery of Heuristic Operators for Vehicle Routing Problems, accessed February 4, 2026, <https://arxiv.org/html/2510.07073v1>
  - 37. Record-Breaking NVIDIA cuOpt Algorithms Deliver Route Optimization Solutions 100x Faster, accessed February 4, 2026,  
<https://developer.nvidia.com/blog/record-breaking-nvidia-cuopt-algorithms-deliver-route-optimization-solutions-100x-faster/>